

Getting started on a Mac

This guide takes you from a Mac with nothing installed to running the two workshop exercises with Codex, OpenAI's coding agent. You do not need any programming experience, because the agent writes and runs the code. On Windows, use the [Windows guide](#) instead.

The guide has seven steps:

1. Download the workshop materials and set up a working folder.
2. Open the Terminal.
3. Install Codex.
4. Install Pixi.
5. Install Git.
6. Run the customer segmentation exercise.
7. Run the Reddit exercise.

If you missed the workshop, the [recording](#) and the [edited transcript](#) show both exercises being started live.

Before you start

Check each item before you install anything.

- **macOS 14 (Sonoma) or newer.** The ChatGPT desktop app needs it. To check, open the Apple menu (top-left corner of the screen) and choose **About This Mac**.
- **At least 10 GB of free disk space.** To check, open the Apple menu, choose **System Settings**, then **General**, then **Storage**.
- **Your Mac login password.** Some installs ask for it. On a Mac managed by your university, installs may be blocked. If they are, ask your IT help desk or use a personal computer. Do not try to work around a university security setting.
- **A ChatGPT account.** The Plus plan (\$20 a month) is the plan these exercises were built for. Eligible U.S. students can get four free months through OpenAI's [Back to School offer](#) by claiming it before October 31, 2026. The offer asks for a payment method and renews at \$20 a month unless you cancel, so put a reminder in your calendar now.
- **If you stay on the Free plan,** Codex works only in the desktop app, with the smaller GPT-6 Luna model. Free access is meant for quick tasks, so you can start the customer exercise, but the later review prompts may not finish.
- **Time and power.** Setup involves several downloads, and the agent works for many minutes after each prompt. Keep the Mac plugged in with the lid open while the agent works, because the work pauses when the Mac sleeps.

1. Download the materials and set up a working folder

You will copy two starter folders, one per exercise, into a new folder named `agentic-ai` in your home folder.

1. Open the [workshop repository](#) in your web browser.
2. Click the green **Code** button, then click **Download ZIP**.
3. Open Finder and press Option+Command+L to open your Downloads folder.
4. Look for a folder named `agentic-ai-fall-2026-manderson-main`. Safari often unzips down-

loads for you. If you see only a file named `agentic-ai-fall-2026-manderson-main.zip`, double-click it to unzip it.

5. In Finder, press `Shift+Command+H` to open your home folder. It has your user name and a house icon.
6. Choose **File**, then **New Folder**. Name the new folder `agentic-ai` and press `Return`.
7. Go back to Downloads and open the `agentic-ai-fall-2026-manderson-main` folder.
8. Click `customer-segmentation-base`, then hold `Command` and click `reddit-base`, so both are selected. Press `Command+C` to copy them.
9. Open the `agentic-ai` folder in your home folder and press `Command+V` to paste.

Check. The `agentic-ai` folder now contains `customer-segmentation-base` and `reddit-base`. Open `customer-segmentation-base`. You should see `SETUP.md`, `transcript.txt`, and a data folder.

Two points about this setup:

- **Why the home folder.** Desktop and Documents are often synced to iCloud. The agent creates thousands of small files, and syncing them causes slowdowns and errors.
- **Why a copy.** If an attempt goes wrong, you can delete the copy and paste a fresh one. The copy also keeps the agent away from the completed examples in the download, so it cannot copy their answers.

2. Open the Terminal

The Terminal is a window where you type commands instead of clicking. You need it to install Pixi and to check that each install worked.

1. Press `Command+Space`, type `Terminal`, and press `Return`.
2. A window opens with a line ending in `%`. This line is the prompt.
3. To run a command, type or paste it after the prompt and press `Return`. Paste with `Command+V`, as usual.

Try it now. Type `cd ~/agentic-ai` and press `Return`, then type `ls` and press `Return`. You should see `customer-segmentation-base` and `reddit-base`.

These commands are all you need:

Command	What it does
<code>pwd</code>	Shows which folder the Terminal is in
<code>ls</code>	Lists the files in that folder
<code>cd agentic-ai</code>	Moves into the folder named <code>agentic-ai</code>
<code>cd ..</code>	Moves up one folder
<code>cd ~</code>	Returns to your home folder (<code>~</code> is short for your home folder)
<code>open .</code>	Opens the current folder in Finder

Three tips:

- Press `Tab` partway through a file or folder name, and the Terminal finishes it for you.

- Type `cd` (with a space after it), then drag a folder from Finder into the Terminal window and press Return. The Terminal moves into that folder.
- Press Control+C to stop a command that is running.

3. Install Codex

Codex comes in two forms. **Use the desktop app** unless you already use the Terminal comfortably. The desktop app is easier to start with and has a button for speaking your prompts. The workshop demo used the command-line version, which runs in the Terminal and needs the Plus plan.

Desktop app (recommended). The desktop app is called ChatGPT, and Codex is one mode inside it.

1. Go to OpenAI's [quickstart page](#) and click the link to download the ChatGPT desktop app.
2. Open the downloaded file and follow its instructions (usually, drag the ChatGPT icon into Applications).
3. Open ChatGPT from Applications and sign in with your ChatGPT account.
4. Choose **Codex** from the menu at the top of the sidebar, which switches between ChatGPT and Codex.

Command-line version (alternative, needs Plus). Paste this into the Terminal and press Return:

```
curl -fsSL https://chatgpt.com/codex/install.sh | sh
```

Then check the install:

1. Close the Terminal window and open a new one. Installers change settings that only new windows read.
2. Type `codex --version` and press Return. You should see a version number.

If you already use Homebrew (a tool for installing Mac software), `brew install --cask codex` works too. The first time you run `codex`, choose **Sign in with ChatGPT** and finish signing in in your browser.

4. Install Pixi

Pixi creates a separate Python environment, i.e., a private set of Python tools, inside each exercise folder. The exercises ask the agent to use it, so you install it once and the agent does the rest.

1. Paste this into the Terminal and press Return:

```
curl -fsSL https://pixi.sh/install.sh | sh
```

2. Close the Terminal window and open a new one.
3. Type `pixi --version` and press Return. You should see a version number.
4. If the ChatGPT app was open during the install, quit it with Command+Q and open it again, so that it can find Pixi.

5. Install Git (recommended)

Git records every change to the files in a folder. With Git installed, Codex can show you exactly what it changed and undo a change. Codex works without Git, but it is easier to follow with it.

1. Type `git --version` in the Terminal and press Return.
2. If you see a version number, Git is already installed. Go to step 6.
3. If a box asks to install the command line developer tools, click **Install** and wait for it to finish. This download can take several minutes.
4. Type `git --version` again. You should see a version number.

6. Run the customer segmentation exercise

The `customer-segmentation-base` folder has three things:

- `transcript.txt`, a dictated description of the assignment.
- `SETUP.md`, instructions written for the agent.
- `data`, eight data files in Parquet, a compact data format. Double-clicking them does nothing useful, and that is expected, because the agent reads them.

Open the folder in Codex

Desktop app.

1. In Codex, add a new project (or open a folder).
2. In the window that opens, press Shift+Command+G, type `~/agentic-ai/customer-segmentation-base`, and press Return. Then click **Open**.
3. Find the permission control beneath the message box and choose **Ask for approval**.
4. Check the model name, which is also beneath the message box. On Plus, choose GPT-6 Sol at medium effort. On Free, GPT-6 Luna is the one available. Effort controls how long the model thinks before it acts.

Command-line version.

1. In the Terminal, type `cd ~/agentic-ai/customer-segmentation-base` and press Return.
2. Type `codex` and press Return.
3. When Codex asks whether to trust the folder, choose the option that lets it work in this folder. Do not choose read-only.
4. The model name appears at the top of the session. Type `/model` to choose GPT-6 Sol at medium effort.

When Codex asks permission

In **Ask for approval** mode, Codex changes files inside the exercise folder on its own. It stops and asks before it uses the internet or touches anything outside the folder.

- **Say yes** when it asks to install packages with `pixi`, to download LaTeX packages, to search Crossref or the web, or to open a file in your browser. The exercises need these. If Codex offers to stop asking about the same kind of command, choosing that option is fine.
- **Say no, and ask it why**, if a command would delete files outside the exercise folder, starts with `sudo`, or asks for your password.

Check your usage before each prompt

Every prompt uses part of your Codex allowance, which resets every five hours and also has a weekly limit. See how much is left on the [usage dashboard](#), or type `/status` in the command-line version. Prompt 4 uses a large share of a Plus plan's weekly allowance, so start it only when most of the week's allowance remains. If Codex offers to sell you extra credits when you reach a limit, you do not need them. Wait for the reset instead.

Send the prompts one at a time

Copy each prompt below into Codex, send it, and read the reply before you send the next one. Working in stages shows you results sooner, and a usage limit is less likely to stop the work halfway.

Prompt 1 asks for questions and a plan before any work begins:

```
Read SETUP.md and transcript.txt. They describe my assignment. Before you run anything, ask me about any choice that would change the result, then give me a step-by-step plan. Use Pixi to create the environment in this folder. If Git is installed, set up a Git repository in this folder so I can see and undo your changes. Do not start the review loop yet.
```

You should see numbered questions and then a plan. No results exist yet. Answer the questions in plain language. It is fine to say, "I don't know, what do you recommend and why?" The agent will likely ask for an email address, because the transcript asks it to search published research through Crossref, a free research database that asks users to identify themselves. Give your university email address.

Prompt 2 runs the analysis and builds the first drafts:

```
The plan looks good. Run the analysis, prototyping on a small sample first, and build first versions of the DOCX report and the HTML explorer. When you finish, summarize what you found, list the exact name and location of every file you created, and open the explorer in my browser.
```

You should see a summary of the customer groups the agent found, a list of files, and the explorer open in your browser. The DOCX report is a Word document, so double-click it in Finder to open it. To open the folder in Finder, type `open .` in the Terminal from the exercise folder, or ask Codex to open it.

Prompt 3 adds one round of independent criticism:

```
Spawn two subagents as independent skeptical reviewers, one focused on methods and one on marketing, as described in SETUP.md. Each should return specific findings with a proposed fix for each. Verify every finding yourself, tell me which ones you accept or reject and why, then apply the accepted fixes.
```

Subagents are extra copies of the agent that Codex starts to review the work. *You should see* a list of findings, each marked accepted or rejected with a reason, and updated files.

Prompt 4 (optional) runs the full scored review loop:

Run the scored review loop described in `SETUP.md`, with three independent reviewer subagents per pass and the scoring and stopping rules it gives. Finish with the process overview.

You should see several rounds of scores that should rise over time, and a process overview, i.e., an HTML page that records the scores and changes.

Understand the result

You are responsible for everything the agent produces for you. Open the report and the explorer, and ask about anything unclear. For example:

- "Why did you choose five segments instead of four?"
- "Walk me through how the spending features are calculated, step by step."

Your results do not need to match the completed example in the customer-segmentation folder of the download. A different, well-justified segmentation is a good outcome.

7. Run the Reddit exercise

The Reddit exercise asks the agent to write a short research paper in LaTeX, a typesetting system common in academic writing. It needs a large data file that is not in the download.

Get the data file

1. Check that `~/agentic-ai/reddit-base` exists. If it does not, copy `reddit-base` from the download as in step 1.
2. Open the [data file on Google Drive](#). You do not need to sign in.
3. Click the download button (a downward arrow, usually near the top right).
4. Google says it cannot scan the file for viruses because it is large. Click **Download anyway**.
5. Wait for the download to finish. The file is about 756 MB.
6. Move `user_daily_post_counts.parquet` from Downloads into `~/agentic-ai/reddit-base/data/`.

Check. In Finder, click the file once and press `Command+I`. The name must be exactly `user_daily_post_counts.parquet`, and the size about 756 MB. If your browser renamed it (for example, with a (1) at the end), rename it.

LaTeX

You do not need to install LaTeX yourself. The first Reddit prompt asks the agent to add Tectonic, a small LaTeX program, inside the exercise folder. The first time the paper compiles, Tectonic downloads the pieces it needs, so say yes when Codex asks.

If you want LaTeX for your own papers later, you can install [MacTeX](#), which is a download of about 6 GB. It is not needed for these exercises.

Send the prompts one at a time

Open `~/agentic-ai/reddit-base` in Codex as in step 6, check your usage, and send these prompts one at a time.

Prompt 1:

```
Read SETUP.md and transcript.txt. They describe my assignment. Before you run anything, ask me about any choice that would change the result, then give me a step-by-step plan. Use Pixi to create the environment in this folder, and add tectonic from conda-forge to that environment to compile the LaTeX paper. If Git is installed, set up a Git repository in this folder. The data file is large, so prototype on a small sample before scanning all of it. Do not start the review loop yet.
```

You should see numbered questions and a plan. As in the customer exercise, give your university email address when it asks for one for Crossref.

Prompt 2:

```
The plan looks good. Run the analysis, prototyping on a small sample first, then on the full file. Write the first complete version of the LaTeX paper and compile it to PDF with tectonic. When you finish, summarize what you found, list the exact name and location of every file you created, and open the PDF.
```

You should see a summary of the findings, a list of files, and the paper open as a PDF. The full data file is large, so this prompt can take a long time.

Prompt 3:

```
Spawn two subagents as independent skeptical reviewers of the paper, with the different perspectives described in SETUP.md. Each should return specific findings with a proposed fix for each. Verify every finding yourself, tell me which ones you accept or reject and why, then apply the accepted fixes and recompile the paper.
```

You should see a list of findings, each marked accepted or rejected with a reason, and an updated PDF.

Prompt 4 (optional):

```
Run the scored review loop described in SETUP.md, with three independent reviewer subagents per round and the scoring and stopping rules it gives. Finish with the process report.
```

You should see several rounds of scores and a process report that records them.

Troubleshooting

Problem	What to do
command not found: codex or command not found: pixi	Close the Terminal window and open a new one. Installers change settings that only new windows read.
The agent says Pixi is not installed	Quit the ChatGPT app with Command+Q and open it again, or open a new Terminal window.
Codex asks for approval before running a command	See "When Codex asks permission" in step 6.
Codex asks before every single file change	You chose read-only. In the app, set the permission control to Ask for approval . In the Terminal, type <code>/permissions</code> and choose the option that lets Codex work in the folder.
An install asks for a password you do not have, or is blocked	Your Mac is probably managed by your university. Ask your IT help desk, or use a personal computer.
A usage limit message appears	Wait for the limit to reset. You do not need to buy credits. Then continue as in the next row.
The Mac went to sleep, or you closed Codex	In the app, click the conversation in the sidebar and send "Continue where you left off." In the Terminal, run <code>codex resume</code> from the exercise folder.
The agent seems stuck	Click the stop button in the app, or press Escape in the Terminal. Then describe what should happen next.
You cannot find a file the agent made	Ask Codex, "Where is the file you just created? Open its folder in Finder."
You want to start an exercise over	Delete the folder in <code>agent1c-ai</code> and copy a fresh one from the download, as in step 1.